

Word Embeddings Tutorial

HILA GONEN

PHD STUDENT AT YOAV GOLDBERG'S LAB

BAR ILAN UNIVERSITY

5/3/18

Outline

- NLP Intro
- Word representations and word embeddings
- Word2vec models
- Visualizing word embeddings
- Word2vec in Hebrew
- Similarity
- Analogies
- Evaluation
- A simple classification example

NLP - Natural Language Processing

NLP is the field that includes

- understanding**
- processing**
- analyzing**
- generating**

natural languages

We aim to create applicative models that perform as similar as possible to humans

NLP

Applications in NLP:

- Translation
- Information Extraction
- Summarization
- Parsing
- Question Answering
- Sentiment Analysis
- Text Classification

And many more...

NLP challenges

This field encounters numerous challenges:

- Polysemy
- Syntactic ambiguity
- Variability
- Co-reference resolution
- Lack of data / huge amounts of data

NLP challenges - Polysemy

Book

Verb: **Book** a flight

Noun: He says it's a very good **book**

Bank

The edge of a river: **He was strolling near the river bank**

A financial institution: **He works at the bank**

Solution

An answer to a problem: **Work out the solution in your head**

From Chemistry: **Heat the solution to 75° Celsius**

NLP challenges – Polysemy

Kids make nutritious snacks

- Kids, when cooked well, can make nutritious snacks

Kids make nutritious snacks

- Kids know how to prepare nutritious snacks

NLP challenges – Syntactic Ambiguity

12 on their way to cruise among dead in plane crash

12 on their way to cruise among dead in plane crash

same words – different meanings

NLP challenges – Syntactic Ambiguity

The cotton clothing is usually made of grows in Mississippi

The cotton clothing is usually made of grows in Mississippi

same words – different meanings

NLP challenges – Syntactic Ambiguity

Fat people eat accumulates

Fat people eat accumulates

same words – different meanings

NLP challenges – Variability

They allowed him to...

They let him ...

He was allowed to...

He was permitted to...

Different words – same meaning

NLP challenges – Co-Reference Resolution

Rachel had to wait for Dan because he said he wanted her advice.

This is a simple case...

There are more complex ones.

Dan called Bob to tell him about his surprising experience last week:
“you won’t believe it, I myself could not believe it”.

NLP challenges – Data-related issues

A lot of data

In some cases, we deal with huge amounts of data

Need to come up with models that can process a lot of data efficiently

Lack of data

Many problems in NLP suffer from lack of data:

- Non-standard platforms (code-switching)
- Expensive annotation (word-sense disambiguation, named-entity recognition)

Need to use methods to overcome this challenge (semi-supervised learning, multi-task learning...)

Representation

We can represent objects in different hierarchy levels:

- Documents
- Sentences
- Phrases
- Words

We want the representation to be interpretable and easy-to-use

Vector representation meets those requirements

We will focus on word representation

The Distributional Hypothesis

The Distributional Hypothesis:

- words that occur in the same contexts tend to have similar meanings
(Harris, 1954)
- “You shall know a word by the company it keeps” (Firth, 1957)

Examples:

- Cucumber, sauce, pizza, ketchup
- Soundtrack, lyrics, sang, duet

tomato

song

Vector Representation

We can define a word by a vector of counts over contexts, For Example:

	song	cucumber	meal	black
tomato	0	6	5	0
book	2	0	2	3
pizza	0	2	4	1

- Each word is associated with a vector of dimension $|V|$ (the size of the vocabulary)
- We expect similar words to have similar vectors
- Given the vectors of two words, we can determine their similarity (more about that later)

We can use different granularities of contexts: documents, sentences, phrases, n-grams

Vector Representation

Raw counts are problematic:

- frequent words will characterize most words -> not informative

Except from raw counts, we can use other functions:

- TF-IDF (for term (t) – document (d)):

$$TF - IDF(t, d) = \frac{count(t, d)}{|d|} \cdot \log \frac{|D|}{|\{d \in D : t \in d\}|} \quad D - \text{set of all documents}$$

- Pointwise Mutual Information:

$$PMI = \log \frac{p(x, y)}{p(x)p(y)}$$

From Sparse to Dense

These vectors are:

- huge – each of dimension $|V|$ (the size of the vocabulary $\sim 100K +$)
- sparse – most entries will be 0

We want our vectors to be small and dense, two options:

1. Use a reduction algorithm such as SVD over a matrix of sparse vectors
2. Learn low-dimensional word vectors directly -
usually referred as “word embeddings”

We will focus on the second option

Word Embeddings

Each word in the vocabulary is represented by a low dimensional vector ($\sim 300d$)

All words are embedded into the same space

Similar words have similar vectors

(= their vectors are close to each other in the vector space)

Word embeddings are successfully used for various NLP applications

Uses of word embeddings

Word embeddings are successfully used for various NLP applications (usually simply for initialization)

- Semantic similarity
- Word sense Disambiguation
- Semantic Role Labeling
- Named entity Recognition
- Summarization
- Question Answering
- Textual Entailment
- Coreference Resolution
- Sentiment analysis
- etc.

Word2Vec

Models for efficiently creating word embeddings

Remember: our assumption is that similar words appear with similar context

Intuition: two words that share similar contexts are associated with vectors that are close to each other in the vector space

Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean, 2013. *Efficient estimation of word representations in vector space*. arXiv preprint arXiv:1301.3781.

Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean, 2013. *Distributed representations of words and phrases and their compositionality*. In Advances in neural information processing systems.

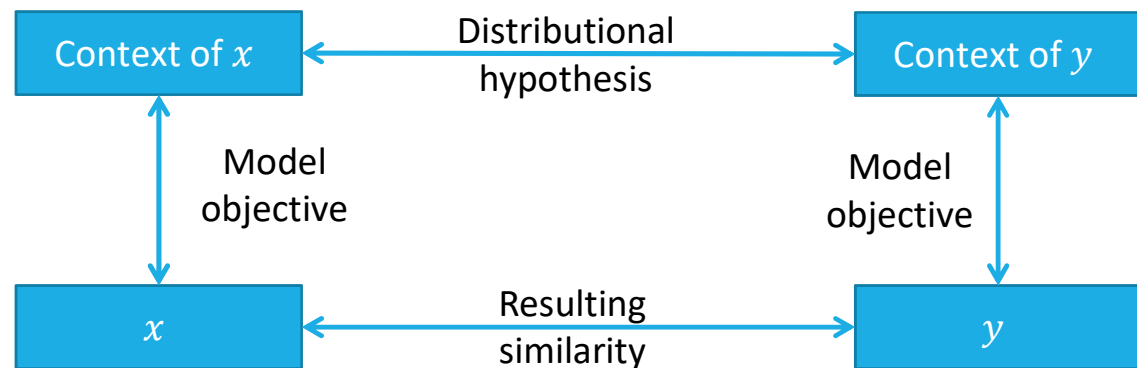
Word2Vec

Models for efficiently creating word embeddings

Remember: our assumption is that similar words appear with similar context

Intuition: two words that share similar contexts are associated with vectors that are close to each other in the vector space

Let x and y be
similar words



Word2Vec

The input: one-hot vectors

- bananas: (1,0,0,0)
- monkey: (0,1,0,0)
- likes: (0,0,1,0)
- every: (0,0,0,1)

vocabulary size
 $|V| = 4$

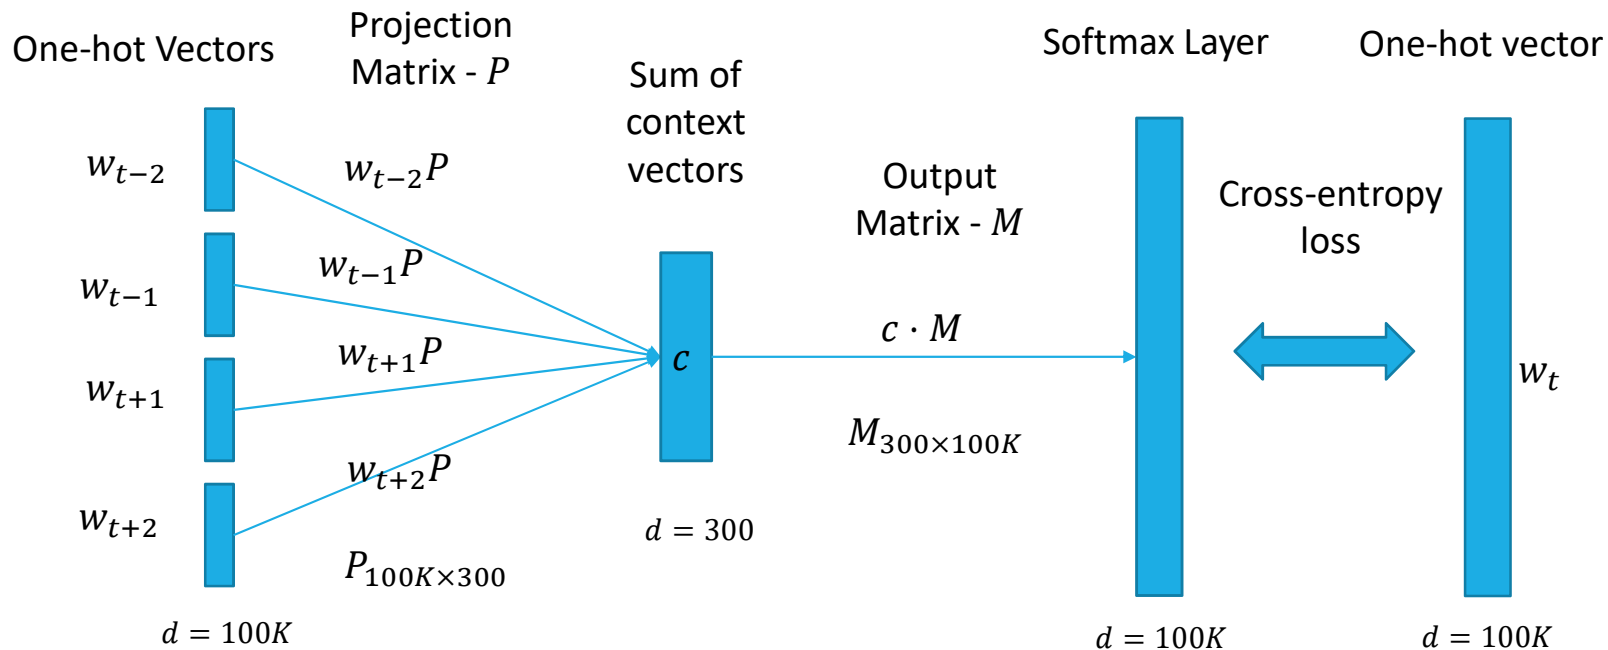
We are going to look at pairs of neighboring words:

Every monkey likes bananas \longrightarrow $(\textit{every}, \textit{monkey})$
 $(\textit{likes}, \textit{monkey})$
 $(\textit{bananas}, \textit{monkey})$

CBOW – high level

The resulting projection matrix P is the embedding matrix


Goal: Predict the middle word given the words of the context



Skip-gram – high level

The resulting projection matrix P is the embedding matrix

Goal: Predict the context words given the middle word



Skip-gram – details

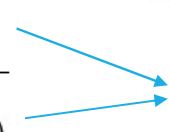
Vector representations will be useful for predicting the surrounding words.

Formally:

Given a sequence of training words w_1, w_2, \dots, w_T , the objective of the Skip-gram model is to maximize the average log probability:

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

The basic Skip-gram formulation defines $p(w_{t+j} | w_t)$ using the softmax function:

$$p(w_{t+j} | w_t) = \frac{\exp(v'_{w_{t+j}} v_{w_t})}{\sum_{i=1}^T \exp(v'_{w_i} v_{w_t})}$$


v - input vector representations
 v' - output vector representations

Negative Sampling

Recall that for Skip-gram we want to maximize the average log probability:

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

Which is equivalent to minimizing the **cross-entropy loss**:

$$L = -\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t) = -\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log \frac{\exp(v'_{w_{t+j}} v_{w_t})}{\sum_{i=1}^T \exp(v'_{w_i} v_{w_t})}$$

This is extremely computational-expensive, as we need to update all the parameters of the model for each training example...

Negative Sampling

When looking at the loss obtained from a single training example, we get:

$$-\log p(w_{t+j}|w_t) = -\log \frac{\exp(v'_{w_{t+j}} v_{w_t})}{\sum_{i=1}^T \exp(v'_{w_i} v_{w_t})} = \underbrace{-v'_{w_{t+j}} v_{w_t}}_{\text{"positive" pair}} + \log \sum_{i=1}^T \exp(\underbrace{v'_{w_i} v_{w_t}}_{\text{"negative" pair}})$$

When using negative sampling, instead of going through all the words in the vocabulary for negative pairs, we sample a modest amount of k words (around 5-20). The exact objective used:

$$\log \sigma(v'_{w_{t+j}} v_{w_t}) + \sum_{i=1}^k \log \sigma(-v'_{w_i} v_{w_t}) \longrightarrow \text{Replaces the term: } \log p(w_{t+j}|w_t) \text{ for each word in the training}$$

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

Context Sampling

We want to give more weight to words closer to our target word

For a given window size C , we sample R in $[1, C]$ and try to predict only R words before and after our target word

For each word in the training we need to perform $2 * R$ word classifications
(R is not fixed)

Subsampling of Frequent Words

In order to eliminate the negative effect of very frequent words such as “in”, “the” etc. (that are usually not informative), a simple subsampling approach is used:

Each word w_i in the training set is discarded with probability: $P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$

This way frequent words are discarded more often

This method improves the training speed and makes the word representations significantly more accurate

Phrases

These models are unable to represent phrases that are not compositions of the individual words

“New York” \neq “New” + “York”

“Boston Globe” \neq “Boston” + “Globe”

The extension is simple:

- Find words that appear frequently together, and infrequently in other contexts
- phrases are formed based on the unigram and bigram counts

$$\text{score}(w_i, w_j) = \frac{\text{count}(w_i w_j) - \delta}{\text{count}(w_i) \times \text{count}(w_j)}$$

- The bigrams with score above the chosen threshold are then used as phrases
- “New York Times” will be replaced with a unique token, “this is” will remain unchanged
- Train word2vec as usual

Use word2vec package

Using this package is extremely simple:

- Download the code from Mikolov's git repository:
- <https://github.com/tmikolov/word2vec>
- Compile the package
- Download the default corpus (wget <http://mattmahoney.net/dc/text8.zip>) or another corpus of your choice
- Train the model using the desired parameters

Jupyter: code for downloading, compiling, and training

Importance of Parameters – window size

Word: **walk**

Window size = 3

Word	Cosine distance
go	0.488083
snipe	0.464912
shoot	0.456677
fly	0.449722
sit	0.449678
pass	0.442459
climbs	0.440931
walked	0.436502
ride	0.434034
stumble	0.426750
bounce	0.425577
travelling	0.419419
walking	0.412107
walks	0.410949
trot	0.410418
leaping	0.406744
sneak	0.401918
climb	0.399793
move	0.396715
wait	0.394463
going	0.391639
shouted	0.388382
roam	0.388073
thrown	0.384087
get	0.383894

Window size = 30

Word	Cosine distance
walking	0.486317
walked	0.430764
walks	0.406772
stairs	0.401518
go	0.399274
sidewalk	0.385786
stand	0.380480
cortege	0.371033
wheelchair	0.362877
strapped	0.360179
hollywood	0.356544
carousel	0.356187
grabs	0.356007
swim	0.355027
breathe	0.354314
tripped	0.352899
cheer	0.352477
moving	0.350943
inductees	0.347791
walkway	0.347164
shout	0.346229
pounding	0.340554
blvd	0.339121
crowd	0.338731
levada	0.334899

Importance of Parameters – iterations

Word: **walk**

No. of iterations = 1

Word	Cosine distance
-----	-----
walking	0.851438
walks	0.846485
bat	0.843796
ride	0.830734
crowd	0.821692
quiet	0.812538
spot	0.802777
steal	0.787917
door	0.787571
doors	0.786485
bed	0.773686
dinner	0.772160
shadow	0.769573
luck	0.768221
baby	0.767862
shoot	0.765968
walked	0.765739
sitting	0.765394
shirt	0.759116
rides	0.759047
watching	0.755140
watch	0.750808
gehrig	0.741494
shoots	0.740971
looking	0.740904

No. of iterations = 100

Word	Cosine distance
-----	-----
walked	0.483473
ride	0.470925
walks	0.470889
stand	0.449993
walking	0.449071
go	0.430172
shoot	0.421110
get	0.404258
move	0.403757
live	0.403347
fly	0.400929
climbs	0.396346
throw	0.391768
climb	0.384038
wiggle	0.380892
thrown	0.380426
pull	0.375478
goes	0.375406
moving	0.374447
pass	0.372463
conversing	0.364413
sit	0.362765
crowd	0.361651
kiss	0.359883
stay	0.357015

Importance of Parameters – dimensions

Word: **walk**

No. of dimensions = 5

Word	Cosine distance
catcher	0.998074
shirt	0.996589
lechuck	0.995313
bullseye	0.994644
bowler	0.994381
punter	0.993154
lovell	0.992815
heels	0.992255
whip	0.992085
outfit	0.992047
tore	0.991924
steals	0.991524
guybrush	0.991166
gigs	0.990291
hanging	0.990201
burns	0.990043
backing	0.989966
orser	0.989960
torch	0.989747
beat	0.989435
showdown	0.989381
feat	0.989242
cheers	0.988951
clad	0.988646
lunch	0.988326

No. of dimensions = 1000

Word	Cosine distance
walks	0.304954
walked	0.303322
snipe	0.287221
walking	0.272690
ride	0.266770
canter	0.251025
bandleaders	0.246454
climbs	0.233725
catapulted	0.230075
climb	0.229263
trot	0.228362
shouted	0.227306
stand	0.223288
seagulls	0.221745
fly	0.216602
fences	0.216366
lifts	0.215402
pray	0.214977
paws	0.214865
bounces	0.214449
shoot	0.213457
grabs	0.212018
walkway	0.211136
swim	0.209120
tumble	0.207765

How does the file usually look like

Word embedding files in readable format usually have a row for each word in the vocabulary

In each row, the specific word is followed by the values of the respected vector

Possibly some additional information in the first rows


Visualizing word embeddings

Using tSNE


- *Visualizing Data using t-SNE*, Maaten and Hinton, 2008

Jupyter: Loading and Visualizing word vectors

Visualizing word embeddings– Word2Vec



Visualizing word embeddings - Glove



Wevi: word embedding visual inspector

A tool that visualizes the basic working mechanism of word2vec

<https://ronxin.github.io/wevi/>

We

A tool

<https://github.com/robertknight/wevi>

wevi: word embedding visual inspector

Everything you need to know about this tool - Source code

Control Panel

Config:

```
{ "hidden_size": 5, "random_state": 1, "learning_rate": 0.2 }
```

Training data (context|target):

```
king|kindom, queen|kindom, king|palace, queen|palace, king|royal, queen|royal, king|George, queen|Mary, man|rice, woman|rice, man|farmer, woman|farmer, man|house, woman|house, man|George, woman|Mary
```

Presets: King and queen

Update and Restart

Update Learning Rate

Next


20

100


500

PCA

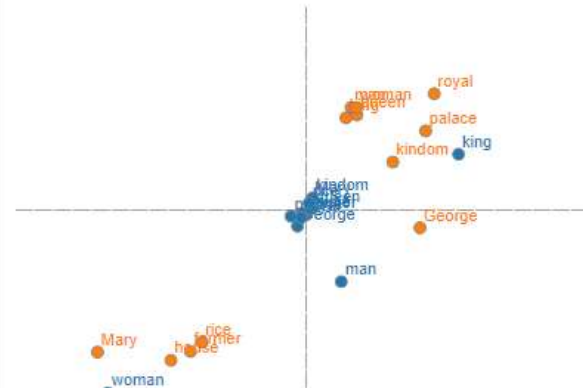
Neurons



Weight Matrices



Vectors



Word2Vec - Hebrew

Link: <http://u.cs.biu.ac.il/~yogo/tw2v/similar/> (By Ron Shemesh and Yoav Goldberg)

Based on tweets in Hebrew

Wor

Link: <http://tw2.net>

Based on

חפש

שכיחות	דימויון מילה
16065	😊 100.0
14738	😊 94.7
20020	😊 93.5
16020	😊 92.7
62507	😊 90.2
68331	😊 87.1
5438	😊 86.6
3556	😊 85.8
4577	😊 82.7
44223	😊 81.5
30838	😊 80.8
19796	😊 79.9
24873	😊 79.6
2327	😊😊 78.8
13376	😊😊 76.5
1236	😊😊 76.7
2197	😊😊 76.6
3758	😊😊 76.0
749	😊😊 75.4
2258	^ ^ 75.3
13669	💜 74.5

ek

/tw2\

חפש

שכיחות	דימויון מילה
327100	הולך 100.0
639232	בא 75.9
680624	יכול 75.3
155222	מתחיל 76.3
166958	הולכת 75.3
87292	אמור 75.3
944253	צריך 71.0
199438	חייב 73.7
1120959	רוצה 68.9
105244	מוכן 73.1
29930	מפסיק 72.8
14106	מחליט 72.7
31218	אלך 72.2
13029	יכל 72.2
5077	ממהר 72.0
28144	שהולך 71.9
60953	ממשיך 70.6
83324	עומד 69.4
37413	מסוגל 69.2
78116	מצליח 68.9
11052	יצטרך 69.1
17427	יתחיל 69.0
177308	מנסה 67.9
47496	רצה 68.1

nesh

חפש

שכיחות	דימויון מילה
20864	שפה 100.0
125	הגיה 78.1
255	הגיה 74.3
288	אסכולה 72.9
292	אינטונציה 71.7
44	דיסציפלינה 71.2
2287	לוגיקה 70.2
261	משפה 70.6
3938	תיאוריה 68.7
496	תפישת 69.3
391	מיתולוגיה 69.2
1099	גנטיקה 68.9
187	ורסיה 69.1
450	כשפה 69.0
179	מתודולוגיה 68.9
3684	פלטפורמה 67.6
589	התחכמות 68.2
9034	הגדרה 66.2
11916	שיטה 65.3
53	סכמה 68.1
1579	פרקטיקה 67.7
582	טרמינולוגיה 67.9
4492	תפיסה 66.9
50	ענה 67.9

Word2Vec - Hebrew

Ok... Nice.

But:

What about יעד vs. מטרה which are synonyms?

Noun genders dramatically affect results –

We do not want that, or at least not for arbitrary gender

Word2Vec

Ok... Nice.

But:

What about יעד vs.

Noun genders drar

We do not want th

חפש יעד

שכיחות	דימיון מילה
8136	יעד 100.0
535	כיעד 70.6
1476	קריטריון 65.2
6366	נכס 63.6
10784	מודל 61.6
158	ויעד 61.6
3626	תרחיש 60.5
11644	היעד 60.0
17600	הישג 59.2
46	כסיכון 59.4
2528	סטנדרט 58.7
38259	פתרון 57.1
1470	תמריץ 56.8
278883	מקום 51.3
37	קריטריון 56.2
95	ובינלאומי 56.0
13	ואקסקלוסיבי 56.0
237	ייתרון 55.8
23446	שטח 55.3
5116	רף 55.1
12953	יתרון 54.6
24049	מהלך 54.1
11696	פוטנציאל 54.2
31019	כיוון 53.9

חפש מטרה

שכיחות	דימיון מילה
19358	מטרה 100.0
2658	שאיפה 74.5
9828	משימה 71.9
61495	הזדמנות 69.1
105895	סיבה 66.2
8805	אלטרנטיבה 70.5
1613	נוסחה 70.7
11916	שיטה 70.0
237	נוסחא 70.5
4714	אסטרטגיה 70.0
7000	אג'נדה 67.9
681	יישות 67.4
496	תפישה 67.4
3684	פלטפורמה 67.0
39317	נקודה 65.1
19288	גישה 65.6
196	דוקטרינה 66.4
44	אסטרטגייה 66.0
21334	בחירה 64.9
4492	תפיסה 65.6
3168	חלופה 65.6
9034	הגדרה 65.3
39166	אפשרות 62.9
1900	ישות 64.6

Word2Vec - Hebrew

Another example:

בצרה

בצל

Prefixes and suffixes are not always handled correctly

Also, not always clear what the wanted behavior is

Word2Vec

Another example

בצרה

בצל

Prefixes and suffixes

Also, not always

חפש בצל

שכיחות	דימיון מילה
23556	בצל 100.0
1574	ובצל 79.8
103	וסלרי 63.3
27207	מרק 56.8
314	ובזיליקום 61.7
286	ופטרוזיליה 61.5
183	מגורד 61.4
5301	ברוטב 59.8
45	חוקיני 60.7
127	בחרדל 60.2
335	וכוסברה 59.9
185	חצתר 60.0
3523	סלק 59.1
7020	רוטב 58.2
1723	ולימון 59.1
298	קונפי 58.9
203	וברוקולי 58.9
196	וסלמון 58.6
77	ואורגנו 58.6
88	ואספרגוס 58.6
147	ופרמזן 58.3
270	כתוש 58.1
235	וסלק 58.1
135	פורטבלו 58.0

ew

is handled

anted behavior

חפש בצרה

שכיחות	דימיון מילה
1624	בצרה 100.0
4622	במצוקה 60.6
8617	בבעיה 50.2
13472	בסביבה 44.5
3026	בצרות 47.2
3730	בהכרה 46.2
9296	בסכנה 42.7
85	ומרעננה 46.1
181	בקרבכם 46.0
184	באחוזה 45.8
97	בקירבה 45.5
16	ונטועה 45.5
12	משמאלם 45.5
1742	בחובות 44.8
96	בטובתם 45.1
1852	במיעוט 44.1
1610	בגלות 44.2
4082	מאחוריך 43.0
2503	בבועה 43.1
27	בבוטניקה 43.7
150	בפרנויה 43.5
775	מאופקים 43.2
53	ובשביה 43.5
61	ובריאותו 43.3

Word2Vec – Hebrew

Word embeddings include inherent biases, as a result of biases in the corpus
(not only in Hebrew...)

For example, רופא vs. חופא

Word2Vec

Word embeddings:

(not only in Hebrew)

For example, רופא

חפש רופאה

שכיחות	דימויון מילה
4307	רופאה 100.0
3064	הרופאה 74.4
2407	מטופלת 73.2
64	פציינטית 72.3
457	וטרנירית 72.1
1242	קוסמטיקאית 71.9
80	עו"סית 71.6
453	כאחות 71.3
2470	פקידה 70.9
3860	תלמידה 69.5
440	מיילדת 69.5
1748	מדריכה 69.1
304	רוקחת 69.1
530	מדענית 68.6
43	קרדיולוגית 68.3
2529	יולדת 67.9
1237	שיננית 67.5
5575	סטודנטית 67.4
4402	גננת 67.2
749	פיליפינית 67.0
329	הוטרינרית 66.7
25801	אחות 66.0
168	פרופסורית 66.5
182017	אישה 61.3

w

biases, a

חפש רופא

שכיחות	דימויון מילה
31223	רופא 100.0
368	אורתופד 85.1
511	כירורג 82.0
3270	פסיכיאטר 79.4
207	נוירולוג 78.0
17732	הרופא 74.6
1772	וטרנר 76.3
218	אורולוג 76.4
5093	פסיכולוג 73.1
298	קרדיולוג 72.6
454	אורטופד 69.8
12820	לרופא 66.1
6271	מטופל 66.7
175	כירורפיקט 66.7
50	סקסולוג 66.6
47	פנימאי 66.3
106	סטאז'ר 66.0
1140	גניקולוג 65.3
309	פתולוג 64.9
52	א.א.ג. 64.4
980	הוטרינר 64.3
4307	רופאה 63.7
450	ורופא 64.0
1630	פסיכיאטרי 63.4

s in the corpus

Other pre-trained word embeddings

Glove (Pennington et al.):

- Based on ratios of co-occurrence probabilities
- <https://nlp.stanford.edu/projects/glove/>

Fast-text (Bojanowski et al.):

- Each word is represented as a bag of character n-grams. A vector representation is associated to each character n-gram, and words are represented as the sum of these representations
- <https://fasttext.cc/>

Similarity

In order to evaluate word embeddings on similarity tasks, we first need to define “similarity”

There are many different ways to define “similarity” and “correlation” between words...

- walk – walking, walk – run, walk – stroll
- Germany – Berlin, Germany – England
- dog – cat, dog – Labrador, dog – leash

This is still an open issue...

Similarity measure

The distance between two vectors is not a good measure

- We do not want to take the length of the vector into account

Most popular similarity measure is Cosine similarity:

- The similarity between two vectors v and w is:
$$\frac{v \cdot w}{||v|| ||w||}$$

Similarity

Jupiter: find topK

Similarity

You can't always get what you want...

walk - Top5 similar words:

walked, walks, walking, climbs, ride

book – top5 similar words:

books, chapter, novel, abridged, autobiography

Analogies

These models are capable of learning linguistic regularities

For example,

$\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) + \text{vector}(\text{"woman"}) \cong \text{vector}(\text{"queen"})$

$\text{vector}(\text{"mice"}) - \text{vector}(\text{"mouse"}) + \text{vector}(\text{"door"}) \cong \text{vector}(\text{"doors"})$

Jupyter: Analogies

Analogies

How does it work?

Given the analogy $a : a^*, b : b^*$, where word b^* is to be found, we try to maximize the following objective:

$$\arg \max_{b^* \in V} (\text{sim}(b^*, b - a + a^*)) \longrightarrow \arg \max_{b^* \in V} (\cos(b^*, b - a + a^*))$$

When vectors are normalized, this is equivalent to:

$$\arg \max_{b^* \in V} (\cos(b^*, b) - \cos(b^*, a) + \cos(b^*, a^*))$$

We actually search for a word that is similar to b , and a^* , but different from a

Linguistic Regularities in Sparse and Explicit Word Representations, Levy and Goldberg, 2014

Analogies

This does not always work that well...

Naturally, depends on the corpus and the hyper-parameters

Additionally, in some cases, specific aspect of relations might dominate others:

London : England , Baghdad : ? \longrightarrow Mosul (instead of Iraq)

Here, even though Iraq is more similar to England than Mosul, the similarity of Mosul to Baghdad dominates, making Mosul the best candidate

A possible solution – use multiplication instead of summation (equivalent to using *log* values):

$$\arg \max_{b^* \in V} \frac{\cos(b^*, b) \cos(b^*, a)}{\cos(b^*, a^*) + \epsilon}$$

Linguistic Regularities in Sparse and Explicit Word Representations, Levy and Goldberg, 2014

Evaluation

Intrinsic Evaluation:

1. Syntactic and semantic analogies:

- Athens : Greece ; Oslo : ?
- think : thinking ; read : ?
- mouse : mice ; door : ?

2. Word correlation benchmarks with human scores (wordsim353, simLex999):

word1	word2	Human score
train	car	6.31
drink	ear	1.31
gem	jewel	8.96

Extrinsic evaluation:

- Show Improvement on downstream tasks when using word embeddings

Classification with word embedding

An optional use of word embedding is a simple classification:

Say we have a short list of professions, and we want to elaborate it

We can run a simple classification model with sklearn

- Use the short list as positive examples
- Add random negative example
- Learn a classification model
- Predict True/False for new words from the vocabulary

Jupyter: Classification example

Biases in word embeddings

A very nice tutorial about word embedding biases:

[How to make a racist AI without really trying](https://gist.github.com/rspeer/ef750e7e407e04894cb3b78a82d66aed)

<https://gist.github.com/rspeer/ef750e7e407e04894cb3b78a82d66aed>

Conclusion

Word embedding:

- A powerful word representation
- Easy to incorporate into different models

Can capture word similarities and linguistic regularities

Existing models have their limitations

Need to custom training parameters according to the desired properties and similarities

Questions?

Word	Cosine distance
question	0.744367
answers	0.577580
doubts	0.535749
answer	0.521965
concerns	0.492096
issues	0.487333
unanswered	0.457177
discussions	0.454697
answering	0.448383
matters	0.446861
debates	0.439423
statements	0.433208
objections	0.429943
issue	0.424728
debate	0.421695
conclusions	0.421594
inquiries	0.419190
problem	0.419025
answered	0.414844
responses	0.412600
arguments	0.410255
problems	0.406869

Thank you!
